

Partial Functions and Preference for ARC: Domains, Interpolation, and Adaptation (Research Note)

Cristiano Calcagno

2025-10-31

Abstract

We formalize Abstraction and Reasoning Corpus (ARC) tasks using partial functions over grids and a numeric preference function on programs. The central concept is the *behavioural domain* that encodes the programmer’s intended coverage. We give a short calculus for behavioural definedness under boolean connectives, characterize interpolation, generalisation, and adaptation to novelty, and work through a representative example where generalisation succeeds precisely by extending the behavioural domain to include a previously unhandled color.

1 Model

Definition 1 (Objects). Let \mathcal{G} be the set of grids. Candidate programs are partial functions $f : \mathcal{G} \rightarrow \mathcal{G}$; write $\text{Dom}_b(f) \subseteq \mathcal{G}$ for the (behavioural) domain of f , i.e. the set of inputs on which f specifies behaviour. Let \mathcal{F} denote a chosen hypothesis class of such programs.

Definition 2 (Preference). A numeric scoring function $P : \mathcal{F} \rightarrow \mathbb{R}$ induces a strict preference: f_1 is preferred to f_2 iff $P(f_1) > P(f_2)$.

Definition 3 (ARC puzzles). An ARC puzzle (instance) consists of training pairs $\{(i_k, o_k)\}_{k=1}^n \subseteq \mathcal{G} \times \mathcal{G}$ and a test input $i \in \mathcal{G}$.

Definition 4 (Interpolants). For training pairs $\{(i_k, o_k)\}_{k=1}^n$, a program f interpolates the training set if for all k , $i_k \in \text{Dom}_b(f)$ and $f(i_k) = o_k$. Let $\mathcal{I} = \{f \in \mathcal{F} \mid f \text{ interpolates } \{(i_k, o_k)\}_{k=1}^n\}$ denote the set of all interpolants.

Definition 5 (Tentative solution, solution, ambiguity/void). A tentative solution is any interpolant $f \in \mathcal{I}$ such that $i \in \text{Dom}_b(f)$. Let

$$\mathcal{S} = \{f \in \mathcal{I} \mid i \in \text{Dom}_b(f)\}.$$

The intended solution is any $f^* \in \arg \max_{f \in \mathcal{S}} P(f)$ (i.e., the element of \mathcal{S} that maximizes P). The instance is void if $\mathcal{S} = \emptyset$ and ambiguous if $|\arg \max_{f \in \mathcal{S}} P(f)| > 1$.

Definition 6 (Generalisation and adaptation to novelty). Generalisation is moving between interpolants $f, g \in \mathcal{I}$ with $P(g) > P(f)$. The top-scoring interpolant $f^\dagger \in \arg \max_{f \in \mathcal{I}} P(f)$ (the interpolant maximizing P) may fail to cover the test input ($i \notin \text{Dom}_b(f^\dagger)$). Adaptation to novelty selects a tentative solution $g^* \in \arg \max_{f \in \mathcal{S}} P(f)$; since $\mathcal{S} \subseteq \mathcal{I}$, necessarily $P(g^*) \leq P(f^\dagger)$.

This formulation clarifies the tension in ARC: generalisation seeks higher-scoring interpolants (simpler, more elegant programs), while adaptation prioritizes coverage of the test input. When the test introduces novel features absent in training (e.g., a new color), the top interpolant f^\dagger may exclude that case from Dom_b , forcing a trade-off: accept a lower-preference tentative solution g^* that extends the behavioural domain, or stick with f^\dagger and risk failure.

2 Behavioural domain

We focus on the **behavioural domain** $\text{Dom}_b(e)$: the inputs on which the specification intends to define behaviour (coverage). We assume standard short-circuit evaluation for expressions; runtime definedness is incidental here and not our emphasis.

Definedness calculus (examples). For boolean expressions E, F and comparisons $x=c$,

$$\text{Dom}_b(\neg E) = \text{Dom}_b(E), \quad (1)$$

$$\text{Dom}_b((x=c \wedge E) \vee F) = (x=c \wedge \text{Dom}_b(E)) \vee \text{Dom}_b(F). \quad (2)$$

For a DSL primitive $g(\cdot)$, we assume a given predicate $\text{Dom}_b(g)$.

3 Worked example

We illustrate the role of the behavioural domain using ARC-AGI-2 task dfadab01. The puzzle is available at <https://arcprize.org/play?task=dfadab01>; the full solver analysis at <https://gist.github.com/cristianoc/e86779cee94658567ecba429133d6667>.

Consider the following (corrected) DSL function:

```

1 def valid_anchor(grid: Grid, anchor: Anchor) -> bool:
2   row, col, color = anchor
3   return not (
4     (color == 2 and guard_nw_eq(grid, (row, col), 4))
5     or (color == 5 and guard_nw_eq(grid, (row, col), 6))
6   )

```

Example 1 (Behavioural domain). *If the intended specification handles only colors 2 (red) and 5 (gray), the behavioural domain must be:*

$$\begin{aligned} \text{Dom}_b(\text{valid_anchor}) = & (color=2 \wedge \text{Dom}_b(\text{guard_nw_eq}(\text{grid}, (\text{row}, \text{col}), 4))) \\ & \vee (color=5 \wedge \text{Dom}_b(\text{guard_nw_eq}(\text{grid}, (\text{row}, \text{col}), 6))). \end{aligned}$$

This captures the explicit color-specific guards: only when color is 2 do we check the guard for color 4, and only when color is 5 do we check the guard for color 6. Any other color falls outside the intended behavioural domain.

Remark 1 (Novelty at test time). *If the test input has $color=8$ (sky blue), then the expression computationally evaluates (no guard is forced), but $i \notin \text{Dom}_b(\text{valid_anchor})$ by the behavioural domain example above. Generalisation must therefore extend the behavioural domain (e.g. add a new branch or a principled fallback) so that 8 becomes covered.*

4 Practical Implications

The model suggests a concrete operating procedure and design tips:

- **Selection rule.** Among interpolants, pick the tentative solutions that cover the test (\mathcal{S}) and choose the P -maximizer; if $\mathcal{S} = \emptyset$ the instance is *void*, and if multiple maximizers remain it is *ambiguous*.
- **Novelty handling.** If the top-scoring interpolant f^\dagger does not cover the test ($i \notin \text{Dom}_b(f^\dagger)$), extend the behavioural domain with the minimal, principled change that includes i , accepting a possible decrease in P .



Figure 1: Before/after comparison showing domain extension failure and recovery: the initial solver explicitly handles only red (2) and gray (5). When sky blue (8) appears in the test input (corresponding to an **X**-pattern motif with orange (7) borders), the solver incorrectly stamps the motif because $8 \notin \text{Dom}_b(\text{valid_anchor})$. The simplified version introduces `guard_color` (deriving the guard color 7 from the motif’s top-left cell) to uniformly treat all motifs, extending Dom_b to include color 8 and preventing spurious stamping.

- **Make coverage explicit.** Encode intended coverage via guard-style case-splits with no default `else`; this makes Dom_b visible and simplifies reasoning about what inputs are intentionally handled.
- **Debugging checklist.** On failures, first distinguish evaluation errors from coverage gaps: compute or approximate $\text{Dom}_b(e)$ for the failing expression and check whether the test input lies outside it.

5 Summary

Modeling ARC programs as partial functions with a numeric score yields a compact account of interpolation, generalisation, and adaptation. The *behavioural domain* is essential: many failures are not evaluation errors but *coverage* gaps. In the example, the initial solver fails on the test because color 8 is outside Dom_b ; the generalised solution succeeds precisely by extending Dom_b to include that case.